

BAYESIAN STATISTICS

ASSIGNMENT 2

QUESTION 1: PROBIT REGRESSION (HOFF 6.3)

A panel study followed $n = 25$ married couples over a period of five years. One item of interest is the relationship between divorce rates and the various characteristics of the couples. For example, the researchers would like to model the probability of divorce as a function of age differential, recorded as the man's age minus the woman's age. The data can be found in the file `divorce.RData`. We will model these data with probit regression, in which a binary variable Y_i is described in terms of an explanatory variable x_i via the following latent variable model:

$$\begin{aligned} Z_i &= \beta x_i + \varepsilon_i \\ Y_i &= \mathbf{1}_{(c, +\infty)}(Z_i), \end{aligned}$$

where β and c are unknown coefficients, $\varepsilon_1, \dots, \varepsilon_n \stackrel{\text{iid}}{\sim} \mathcal{N}(0, 1)$ and $\mathbf{1}_{(c, +\infty)}(z) = 1$ if $z > c$ and equals zero otherwise. In the following, since the covariates x_i are known, they will be treated as constants and so not explicitly written in the conditioning part.

Point a.

Assuming $\beta \sim \mathcal{N}(0, \sigma_\beta^2)$, obtain the full conditional distribution $p(\beta | y_{1:n}, z_{1:n}, c)$.

First of all let us write explicitly the conditional distributions which we can deduce from the text:

– $\forall i = 1, \dots, n$ we know $p(z_i | \beta)$:

$$\begin{aligned} Z_i(\omega) | \beta &= \beta x_i + \varepsilon_i(\omega) \sim \beta x_i + \mathcal{N}(0, 1) \sim \mathcal{N}(\beta x_i, 1) \implies Z_i | \beta \sim \mathcal{N}(\beta x_i, 1) \\ &\Downarrow \\ p(z_i | \beta) &= \frac{1}{\sqrt{2\pi}} e^{-\frac{1}{2}(z_i - \beta x_i)^2}; \end{aligned}$$

– $\forall i = 1, \dots, n$ we know $p(y_i | c, z_i)$:

$$\begin{aligned} Y_i(\omega) &= \mathbf{1}_{(c, +\infty)}(Z_i) = \begin{cases} 1 & \text{if } Z_i > c \\ 0 & \text{otherwise} \end{cases} \\ &\Downarrow \\ p(y_i) &= \mathbb{P}(Y_i = y_i) = \mathbb{P}(\mathbf{1}_{(c, +\infty)}(Z_i) = y_i) = \\ &= \begin{cases} \mathbb{P}(\mathbf{1}_{(c, +\infty)}(Z_i) = 1) & \text{if } y_i = 1 \\ \mathbb{P}(\mathbf{1}_{(c, +\infty)}(Z_i) = 0) & \text{if } y_i = 0 \\ 0 & \text{otherwise} \end{cases} = \\ &= \begin{cases} \mathbb{P}(\{Z_i > c\}) & \text{if } y_i = 1 \\ \mathbb{P}(\{Z_i > c\}^C) & \text{if } y_i = 0 \\ 0 & \text{otherwise} \end{cases} = \\ &= (y_i \mathbb{P}(\{Z_i > c\}) + (1 - y_i) \mathbb{P}(\{Z_i > c\}^C)) \mathbf{1}_{\{0,1\}}(y_i), \end{aligned}$$

hence $Y_i \sim \text{Bernoulli}(\mathbb{P}(Z_i > c))$.

It follows that, conditionally on Z_i, c , the r.v. Y_i is no more *random* and it holds¹

$$p(y_i | c, z_i) = \left(y_i \mathbb{1}_{(-\infty, z_i)}(c) + (1 - y_i) \mathbb{1}_{(-\infty, z_i)^c}(c) \right) \mathbb{1}_{\{0,1\}}(y_i).$$

In order to obtain (and sample) from the full conditionals we assume β and c a priori independent. The full conditional distribution $p(\beta | y_{1:n}, z_{1:n}, c)$ can be obtained just from $p(z_i | \beta)$, indeed

$$\begin{aligned} p(\beta | y_{1:n}, z_{1:n}, c) &= \frac{p(\beta, y_{1:n}, z_{1:n}, c)}{p(y_{1:n}, z_{1:n}, c)} \frac{p(\beta, z_{1:n}, c)}{p(\beta, z_{1:n}, c)} \frac{p(\beta, c)}{p(\beta, c)} \frac{p(c)}{p(c)} \propto \\ &\propto \frac{p(\beta, y_{1:n}, z_{1:n}, c)}{p(\beta, z_{1:n}, c)} \frac{p(\beta, z_{1:n}, c)}{p(\beta, c)} \frac{p(\beta, c)}{p(c)} = \\ &= p(y_{1:n} | \beta, c, z_{1:n}) p(z_{1:n} | \beta, c) p(\beta | c) \propto \\ &\propto p(z_{1:n} | \beta) p(\beta). \end{aligned}$$

So we can write explicitly

$$\begin{aligned} p(\beta | y_{1:n}, z_{1:n}, c) &\propto p(z_{1:n} | \beta) p(\beta) = \\ &= \prod_{i=1}^n p(z_i | \beta) p(\beta) \propto \\ &\propto \exp \left(-\frac{1}{2} \sum_{i=1}^n (z_i - x_i \beta)^2 \right) \exp \left(-\frac{1}{2} \frac{1}{\sigma_\beta^2} \beta^2 \right) = \\ &= \exp \left(-\frac{1}{2} \left(\beta \sum_{i=1}^n x_i^2 + \sum_{i=1}^n z_i^2 - 2\beta \sum_{i=1}^n x_i z_i + \beta^2 \frac{1}{\sigma_\beta^2} \right) \right) = \\ &= \exp \left(- \underbrace{\left(\sum_{i=1}^n x_i^2 + \frac{1}{\sigma_\beta^2} \right)}_{\stackrel{\text{def}}{=} (\sigma_{\beta,n}^2)^{-1}} \frac{\beta^2}{2} + \underbrace{\left(\sum_{i=1}^n x_i z_i \right)}_{\stackrel{\text{def}}{=} \frac{\mu_{\beta,n}}{\sigma_{\beta,n}^2}} \beta \right), \end{aligned}$$

where from the 1st to the 2nd line we used $(Z_i | \beta)_{i=1}^n$ independent, identically distributed r.v.'s. So we can conclude that

$$\begin{aligned} \beta | y_{1:n}, z_{1:n}, c &\sim \mathcal{N} \left(\mu_{\beta,n}, \sigma_{\beta,n}^2 \right) \text{ with } \begin{cases} \sigma_{\beta,n}^2 = \left(\sum_{i=1}^n x_i^2 + \frac{1}{\sigma_\beta^2} \right)^{-1} \\ \mu_{\beta,n} = \sigma_{\beta,n}^2 \left(\sum_{i=1}^n x_i z_i \right) \end{cases} \\ &\Downarrow \\ p(\beta | y_{1:n}, z_{1:n}, c) &= \frac{1}{\sqrt{2\pi\sigma_{\beta,n}^2}} \exp \left(-\frac{1}{2\sigma_{\beta,n}^2} (\beta - \mu_{\beta,n})^2 \right). \end{aligned}$$

□

Point b.

Assuming $c \sim \mathcal{N}(0, \sigma_c^2)$, show that $p(c | y_{1:n}, z_{1:n}, \beta)$ is a constrained normal density, i.e. proportional to a normal density but constrained to lie in an interval. Similarly, show that $p(z_i | y_{1:n}, z_{-i}, \beta, c)$ is proportional to a normal density but constrained to be either above c or below c , depending on y_i .

¹We replace $\mathbb{P}(\{z_i > c\})$ with $\mathbb{1}_{(-\infty, z_i)}(c)$ because we will use this characterization afterwards.

Hint: A constrained, or truncated, normal random variable V is obtained by restricting a normally distributed random variable $\mathcal{N}(\mu, \tau^2)$ to lie in an interval (a, b) , with possibly $a = -\infty$ or $b = +\infty$. We use the notation $V \sim \mathcal{TN}_{(a,b)}(\mu, \tau^2)$. It holds:

- $p(v | \mu, \tau^2, a, b) = \frac{1}{C} \frac{1}{\sqrt{2\pi\tau^2}} \exp\left(-\frac{1}{2\tau^2}(v - \mu)^2\right) \mathbb{1}_{(a,b)}(v)$, where $C = \Phi\left(\frac{b-\mu}{\tau}\right) - \Phi\left(\frac{a-\mu}{\tau}\right)$ being $\Phi(\cdot)$ the cdf of the standard normal distribution. By definition, it holds $\Phi\left(\frac{b-\mu}{\tau}\right) = 1$ if $b = \infty$ and $\Phi\left(\frac{a-\mu}{\tau}\right) = 0$ if $a = -\infty$.
- Sampling can be performed thanks to the function `rtruncnorm(n, a, b, mean, sd)` from the package `rtruncnorm` [<https://cran.r-project.org/web/packages/truncnorm/truncnorm.pdf>]. This function receives in input the number of desired samples (n) and the four parameters specifying the distribution of V : a, b, μ, τ . Pay attention that it takes as last inputs the mean μ and the standard deviation τ (not the variance τ^2) of the un-truncated normal density.

As before, the full conditional distribution $p(c | y_{1:n}, z_{1:n}, \beta)$ can be obtained just from $p(y_i | c, z_i)$, indeed

$$\begin{aligned} p(c | y_{1:n}, z_{1:n}, \beta) &= \frac{p(c, y_{1:n}, z_{1:n}, \beta)}{p(y_{1:n}, z_{1:n}, \beta)} \frac{p(\beta, c, z_{1:n})}{p(\beta, c, z_{1:n})} \frac{p(c, \beta)}{p(c, \beta)} \frac{p(\beta)}{p(\beta)} \propto \\ &\propto \frac{p(c, y_{1:n}, z_{1:n}, \beta)}{p(\beta, c, z_{1:n})} \frac{p(\beta, c, z_{1:n})}{p(c, \beta)} \frac{p(c, \beta)}{p(\beta)} = \\ &= p(y_{1:n} | \beta, c, z_{1:n}) p(z_{1:n} | \beta, c) p(c | \beta) \propto \\ &\propto p(y_{1:n} | c, z_{1:n}) p(c). \end{aligned}$$

So we can write explicitly

$$\begin{aligned} p(c | y_{1:n}, z_{1:n}, \beta) &\propto p(y_{1:n} | c, z_{1:n}) p(c) = \\ &= \prod_{i=1}^n p(y_i | c, z_i) p(c) \propto \\ &\propto \exp\left(-\frac{1}{2} \frac{1}{\sigma_c^2} c^2\right) \prod_{i=1}^n \left(y_i \mathbb{1}_{(-\infty, z_i)}(c) + (1 - y_i) \mathbb{1}_{(-\infty, z_i)^c}(c)\right) \mathbb{1}_{\{0,1\}}(y_i) = \\ &= \exp\left(-\frac{1}{2} \frac{1}{\sigma_c^2} c^2\right) \prod_{i=1, \dots, n | y_i=1} \mathbb{1}_{(-\infty, z_i)}(c) \cdot \prod_{i=1, \dots, n | y_i=0} \mathbb{1}_{[z_i, +\infty)}(c) = \\ &= \exp\left(-\frac{1}{2} \frac{1}{\sigma_c^2} c^2\right) \mathbb{1}_{(-\infty, \min(z_i | i \in \{1, \dots, n\}, y_i=1))}(c) \mathbb{1}_{[\max(z_i | i \in \{1, \dots, n\}, y_i=0), +\infty)}(c), \end{aligned}$$

where from the 1st to the 2nd line we used $(Y_i | c, z_i)_{i=1}^n$ independent, identically distributed r.v.'s. More compactly, defining

$$\begin{aligned} a_n &\stackrel{\text{def}}{=} \max(z_i | i \in \{1, \dots, n\}, y_i = 0) \text{ and} \\ b_n &\stackrel{\text{def}}{=} \min(z_i | i \in \{1, \dots, n\}, y_i = 1), \end{aligned}$$

one has

$$p(c | y_{1:n}, z_{1:n}, \beta) \propto \exp\left(-\frac{1}{2} \frac{1}{\sigma_c^2} c^2\right) \mathbb{1}_{[a_n, b_n)}(c).$$

First of all we have to observe that the indicator function constrains $c \in [a_n, b_n)$, but it is equivalent to $c \in (a_n, b_n)$ because our $p(c | y_{1:n}, z_{1:n}, \beta)$ is a density function with respect to the lebesgue measure on \mathbb{R} so each point has measure 0 (so does $\{a_n\}$).

Then, let us observe that this conditional density is proportional to the kernel of a gaussian (evaluated in c) multiplied by an indicator function (also evaluated in c), which constrains the domain to an interval

(not necessarily limited, possibly $a_n = -\infty$ or $b_n = +\infty$).

So completing the function $\exp\left(-\frac{1}{2}\frac{1}{\sigma_c^2}c^2\right)\mathbb{1}_{(a_n,b_n)}(c)$ to a density one obtains

$$\begin{aligned} p(c | y_{1:n}, z_{1:n}, \beta) &= \frac{1}{\Phi\left(\frac{b_n}{\sigma_c}\right) - \Phi\left(\frac{a_n}{\sigma_c}\right)} \frac{1}{\sqrt{2\pi\sigma_c^2}} \exp\left(-\frac{1}{2}\frac{1}{\sigma_c^2}c^2\right) \mathbb{1}_{(a_n,b_n)}(c) \\ &\Downarrow \\ c | y_{1:n}, z_{1:n}, \beta &\sim \mathcal{TN}_{(a_n,b_n)}\left(0, \sigma_c^2\right). \end{aligned}$$

Similarly

$$\begin{aligned} p(z_i | y_{1:n}, z_{-i}, \beta, c) &= \frac{p(z_i, y_{1:n}, z_{-i}, \beta, c)}{p(y_{1:n}, z_{-i}, \beta, c)} \frac{p(z_i, z_{-i}, \beta, c)}{p(z_i, z_{-i}, \beta, c)} \frac{p(z_i, \beta, c)}{p(z_i, \beta, c)} \frac{p(\beta, c)}{p(\beta, c)} \propto \\ &\propto \frac{p(z_i, y_{1:n}, z_{-i}, \beta, c)}{p(z_i, z_{-i}, \beta, c)} \frac{p(z_i, z_{-i}, \beta, c)}{p(z_i, \beta, c)} \frac{p(z_i, \beta, c)}{p(\beta, c)} = \\ &= p(y_{1:n} | z_{1:n}, \beta, c) p(z_{-i} | z_{-i}, \beta, c) p(z_i | \beta, c) \propto \\ &\propto p(y_{1:n} | z_{1:n}, c) p(z_i | \beta) \propto \\ &\propto \prod_{j=1}^n p(y_j | z_j, c) p(z_i | \beta) \propto \\ &\propto p(y_i | z_i, c) p(z_i | \beta) \propto \\ &\propto \left(y_i \underbrace{\mathbb{1}_{(-\infty, z_i)}(c)}_{= \mathbb{1}_{(c, +\infty)}(z_i)} + (1 - y_i) \underbrace{\mathbb{1}_{(-\infty, z_i)}(c)}_{= \mathbb{1}_{(-\infty, c]}(z_i)} \right) \mathbb{1}_{\{0,1\}}(y_i) \exp\left(-\frac{1}{2}(z_i - x_i\beta)^2\right) = \\ &= \begin{cases} \mathbb{1}_{(c, +\infty)}(z_i) \exp\left(-\frac{1}{2}(z_i - x_i\beta)^2\right) & \text{if } y_i = 1 \\ \mathbb{1}_{(-\infty, c]}(z_i) \exp\left(-\frac{1}{2}(z_i - x_i\beta)^2\right) & \text{if } y_i = 0 \end{cases}. \end{aligned}$$

As before, this conditional density is proportional to the kernel of a gaussian (evaluated in z_i) multiplied by an indicator function (also evaluated in z_i) which constrains the domain to be $(c, +\infty)$ or $(-\infty, c]$ (equivalently $(-\infty, c)$, with the same motivation given above) depending on y_i .

In particular, completing to a density what we found

$$\begin{aligned} p(z_i | y_{1:n}, z_{-i}, \beta, c) &= \begin{cases} \frac{1}{1 - \Phi(c - x_i\beta)} \frac{1}{\sqrt{2\pi}} \exp\left(-\frac{1}{2}(z_i - x_i\beta)^2\right) \mathbb{1}_{(c, +\infty)}(z_i) & \text{if } y_i = 1 \\ \frac{1}{\Phi(c - x_i\beta)} \frac{1}{\sqrt{2\pi}} \exp\left(-\frac{1}{2}(z_i - x_i\beta)^2\right) \mathbb{1}_{(-\infty, c)}(z_i) & \text{if } y_i = 0 \end{cases} \\ &\Downarrow \\ Z_i | y_{1:n}, z_{-i}, \beta, c &\sim \begin{cases} \mathcal{TN}_{(c, +\infty)}(x_i\beta, 1) & \text{if } y_i = 1 \\ \mathcal{TN}_{(-\infty, c)}(x_i\beta, 1) & \text{if } y_i = 0 \end{cases}. \end{aligned}$$

□

Point c.

Letting $\sigma_\beta^2 = \sigma_c^2 = 16$, implement a Gibbs sampling scheme that approximates the joint posterior distribution of $Z_{1:n}, \beta$ and c . After a burnin of 1000, run the Gibbs sampler long enough so that the effective sample sizes of all unknown parameters are greater than 1000 (including the Z_i 's). Compute the autocorrelation function of the parameters and discuss the mixing of the Markov chain.

Point d.

Obtain a 95% posterior credible interval for β , as well as $\mathbb{P}(\beta > 0 \mid y_{1:n})$.